



Intermediate HPC

Introduction to Unix for HPC

What is HPC?

- HPC, or high-performance computing, refers to the application of supercomputers or clusters of computers to computational problems that typically arise through scientific inquiry
- HPC is useful when a computational problem:
 - **Is too large** to solve on a conventional laptop or workstation (because it requires too much memory or disk space) or
 - **Would take too long** (because the algorithm is complex, the dataset is large, or data access is slow) or
 - **Are too many** – High Throughput Computing

Parallelism on HPC

- HPC systems often derive their computational power by exploiting parallelism
- Programs for HPC systems must be split up into many smaller “sub-programs” which can be executed in parallel on different processors
- HPC systems can offer parallelism at a much larger scale, with 100’s or 1000’s, or (soon) even millions of tasks running concurrently.
- Writing parallel software can be challenging, and many existing software packages do not already support parallelism & may require development.

NOTE: Many tasks cannot be parallelised

Reasons to use HPC

- You have a program that can be recompiled or reconfigured to use optimized numerical libraries that are available on HPC systems but not on your own system.
- HPC applications are already installed on the HPC machines which is a non-trivial task
- You have a "parallel" problem, e.g. you have a single application that needs to be rerun many times with different parameters.
- You have an application that has already been designed with parallelism
- To make use of the large memory available
- Our facilities are reliable and regularly backed up

When not to use HPC?

- You have a single threaded job which will only run one job at a time (typical of MatLab users)
- You rely on Databases
- You have a lot of data to transfer between your local machine and the HPC on a continuous basis (e.g. per job)
- You need to have a GUI to interact with your program

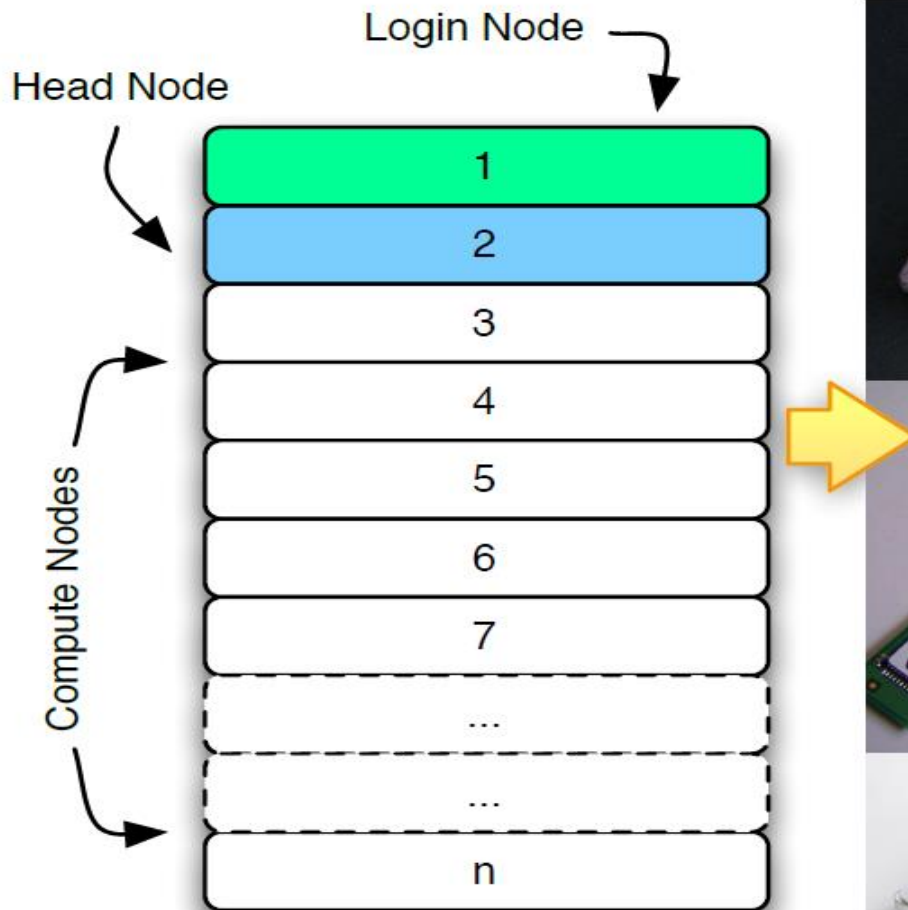
HPC machines

| System | Memory Architecture | Cores | Nodes | Memory |
|---------------------------|---------------------|---------------|-------------|--------------|
| Octane (training machine) | Distributed | 48 | 3 | 48GB |
| Orange | Distributed | 1,600 | 100 | 8TB |
| NCI – (Raijin) | Distributed | 57,472 | 3592 | 158TB |


The typical HPC workflow

- In HPC we talk about **jobs**, these are simply commands we wish to run and requests for resources (e.g. compute time, disk space, memory requirements, setup of s/w env's etc.)
- Generally time consuming & resource intensive.
- Jobs are typically run **non-interactively**,
- Can be run **interactively** for testing purposes
- We add our jobs to a **queue**.
- When machines have **free resources** jobs run
- Once jobs are complete, we can inspect their output.

The HPC "Cluster"



Each node



Processors




Memory



Disk


The Login Node



Login Node

- Interactive programs
- SSH sessions
- Testing
- Compiling


The Head Node



Head Node

- Queuing jobs

Compute Nodes



- These nodes run your jobs
- Managed by the **scheduler**
- Typically you won't interact with the nodes directly
- Some users may need to!

Queuing Systems

- **Portable Batch System** (PBS) is the name of computer software that performs job scheduling. Its primary task is to allocate computational tasks, i.e., batch jobs, among the available computing resources.
- The following versions of PBS are currently available:
 - OpenPBS
 - TORQUE
 - PBS Professional (PBS Pro)
 - ANU PBS
- Guide to PBS: <http://hpc.sissa.it/pbs/pbs.html>

Queuing Systems cont.

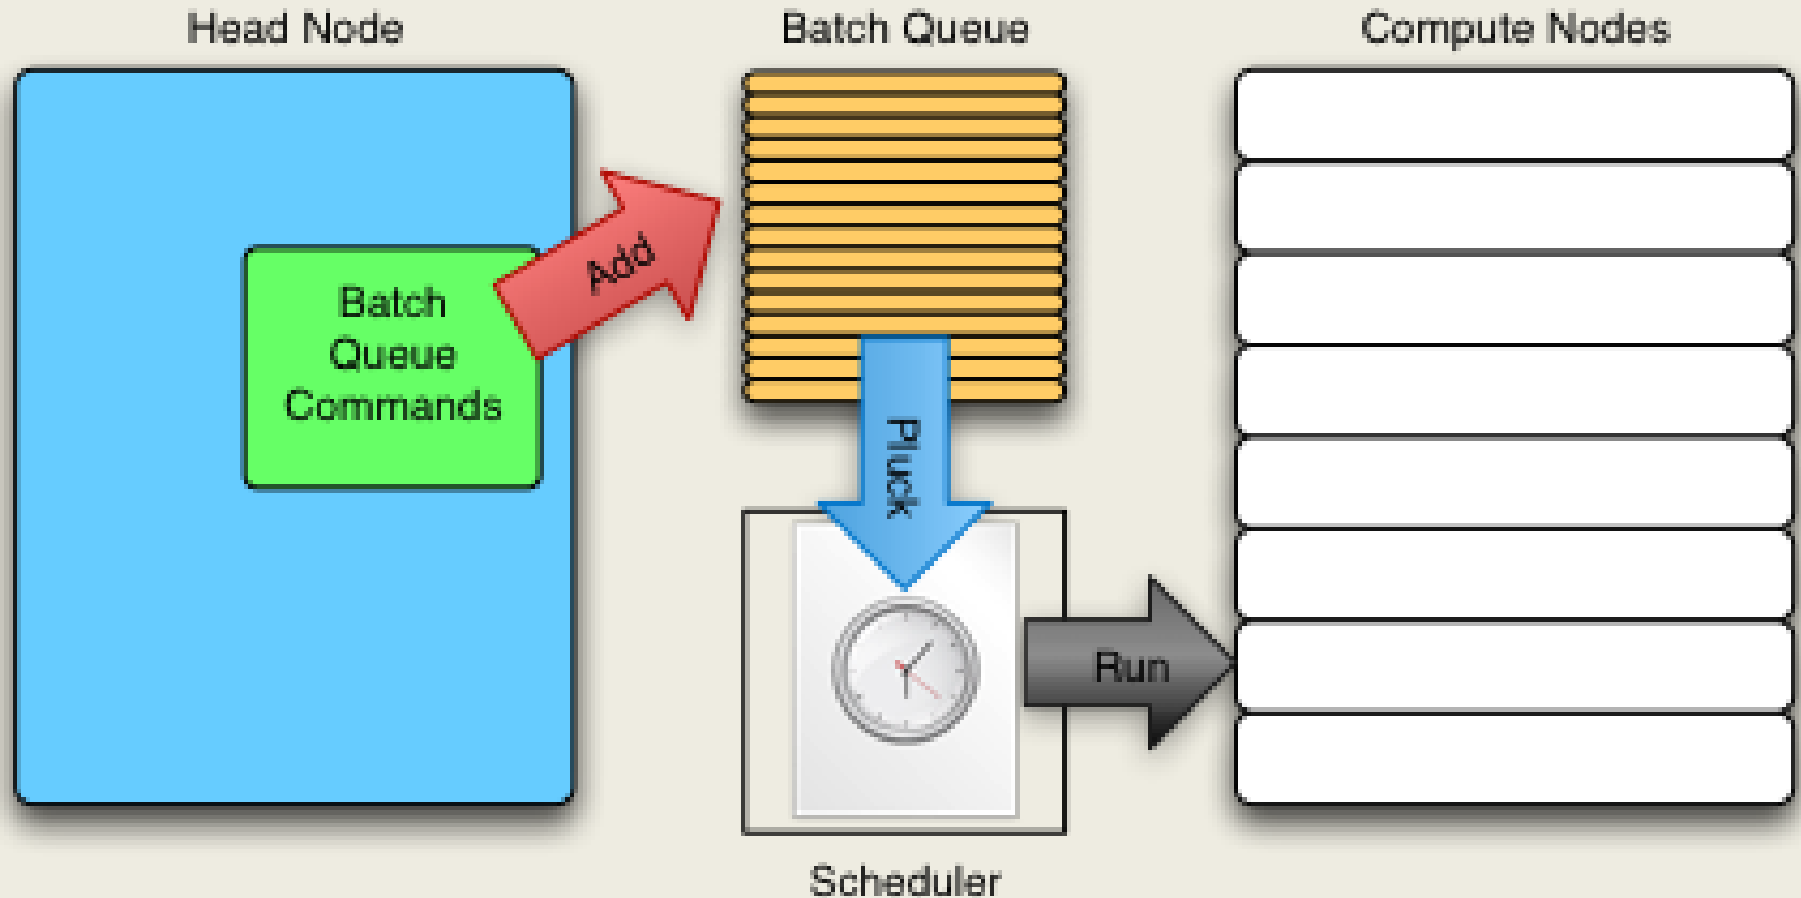
- Another popular batch system is **SLURM** (Simple Linux Utility for Resource Management)
 - Open source, fault-tolerant, and highly scalable cluster management and job scheduling system for large and small Linux clusters.
 - Very useful for use on clusters
 - Platform Tools used by IBM
 - Used by many supercomputers, e.g. TERA 100 at CEA (Europe's most powerful supercomp.)
- Many banks and commercial entities using batch systems

ANU PBS vs PBS Pro

- ANU PBS is a **customised version of PBS based on OpenPBS 2.3** maintained by ANU
- Details of ANU PBS modifications are found here:
http://anusf.anu.edu.au/~dbs900/PBS/local_modifications.html


| Batch System | ANU PBS | PBS PRO |
|----------------|-------------|------------------|
| Machines using | Raijin | Orange & Octane |
| Code Base | OpenPBS 2.3 | PBS Professional |
| Licence | ANU Licence | Altair Licence |

The Batch Queuing System



Batch Queuing component

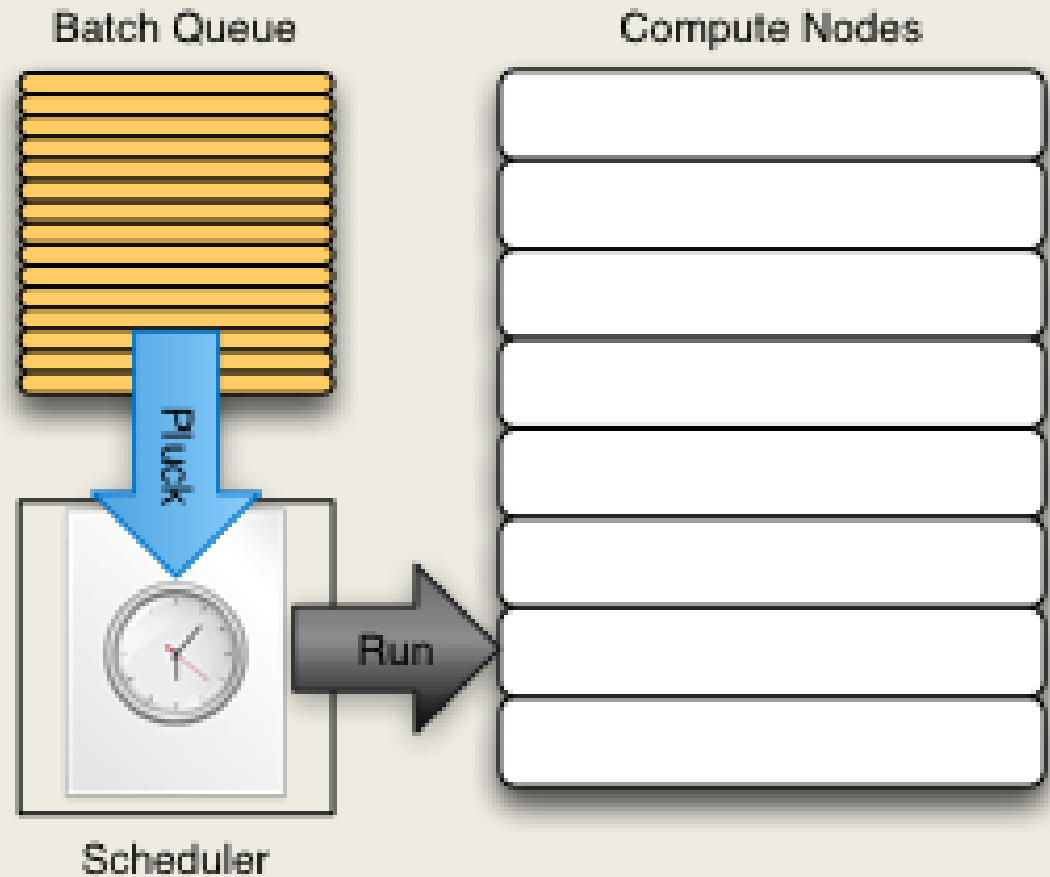
Head Node



- The batch system is a normal program
- Lets you add and remove jobs from the queue and monitor the queue
- Script/command line driven

The Scheduler component

- Allocates jobs to compute nodes
- Optimizes usage of resources
- “Optimize” can mean many things
- Non-trivial
- Never interact with directly



PBS Pro Commands

In order to use the batch system productively, we need to know how to perform three actions:

- Add a job to the queue
- Remove a job from the queue
- See where our job is in the queue

| Command | Description |
|--|---|
| <code>qsub <job-script></code> | Submit a job (add to queue) Returns a <job-number> |
| <code>qdel <job-number></code> | Delete job (remove from the queue) |
| <code>qstat <job-number></code> | Monitor jobs |
| <code>qalter <job-number></code> | Modifies the attributes of the job or jobs |

Exercise 1

Monitoring the queue with qstat

| Command | Description |
|--|--|
| <code>qstat -a</code> | List all jobs in the queue |
| <code>qstat -u <username></code> | List all jobs of a particular user |
| <code>qstat -f <job-number></code> | Show detailed information about a job |
| <code>nqstat</code> | ANU PBS only! Shows detailed information about all jobs |

All about Modules

- **What do Modules do?**
 - Set up the environment for a software package
 - Adds paths to executables to \$PATH
 - May change other shell variables, and/or load other modules.
 - Allow you to have different versions of the same software package, e.g. load the Intel compilers v23 with module load intel/12, or if you require an older version, you can load it with module load intel/9

All about Modules (cont.)

Things to know about modules

- No default module loaded when you login. You can change this by adding lines to your .bashrc file (BASH users) or similar for other Shells
- Modules exist for many packages but not for all!
- Some modules exclude each other, e.g.
 - You can't load Intel compilers v8 & v9. You can only load one.
 - You can only load 1 MPI (e.g. MPT or Intel or OpenMPI)

Module Commands

| Command | Description |
|---|--|
| <code>module avail</code> | Will list all available module files in the current MODULEPATH |
| <code>module load/unload</code> | Will load/unload a modulefile into the shell environment e.g. <code>module load mpt/2.06</code> |
| <code>module list</code> | List all loaded modules |
| <code>module show [modulefile]</code> | Will show information about the modulefile |

Add a job to the queue

- To add a job to the queue, we write a **job script**. A job script is simply a script.
- The **#** symbol signifies a comment for the Shell
- It has some special comments that pass info to PBS Pro.
- **#PBS** is a keyword for PBS and specifies that this line is for PBS. The Shell will ignore it
- When we want to queue the job, we pass its filename as a parameter to **qsub**, e.g.
 - **qsub <job-script>**
- The batch queuing system will return a number that uniquely identifies the job.

Useful Environment Variables

- These are available in the context of your job script.

| Command | Description |
|----------------------------|---|
| <code>PBS_O_WORKDIR</code> | The directory the job was submitted from |
| <code>PBS_JOBID</code> | The job number given when the job was submitted |

A sample PBS job script

NOTE: This script will not run on Orange as resources are assigned via Groups on Orange, not via projects

```
#!/bin/bash
# Request resources
# * 10 minutes wall time to run
#PBS -l walltime=00:10:00
# * 1 node, 1 processor
#PBS -l nodes=1:ppn=1
# * 100 megabytes physical memory allocated to job
#PBS -l mem=100mb
# Specify a project code (for accounting)
#PBS -P a40
cd $PBS_O_WORKDIR
# Specify the job to be done
date
sleep 10
date
```

You've got mail!

```
# Set email address
```

```
#PBS -M fred@intersect.org.au
```

```
# Send an email when jobs
```

```
# begins (b), gets aborted (a)
```

```
# and ends (e)
```

```
#PBS -m abe
```

Exercise 2

Create a script and submit a very simple job

| Command | Description |
|-------------------------------------|---|
| #PBS | #PBS is a keyword for PBS and specifies that this line is for PBS. The Shell will ignore it |
| # | Signifies a comment for the Shell, e.g. # Next line will create a job |
| qsub | Submit a job to PBS |
| qstat | List jobs in the queue |
| cat <i><filename></i> | Print a file to the terminal (catenate) |
| less <i><filename></i> | Like cat , but less at a time |
| nano <i><filename></i> | Will open file <i><filename></i> in a text file editor |
| Sample PBS Script | http://www.intersect.org.au/orange-handbook |

Exercise 3

Create another script and submit a more realistic sample job

| Command | Description |
|-------------------------------------|--|
| #PBS | #PBS is a keyword for PBS and specifies that this line is for PBS. The Shell will ignore it |
| # | Signifies a comment for the Shell, e.g. # Next line will create a job |
| qsub | Submit a job to PBS |
| qstat | List jobs in the queue |
| nano <filename> | Will open file <filename> in a text file editor |
| module load <module_name> | Loads a software module on the HPC machine |

Job limits on Orange

- 200 hours of **walltime**
- 64GB of **memory** per standard node, e.g. 128GB for 2 nodes etc.
- 256GB of **memory** per large memory nodes
- **NOTE**: If you grab a node with 64GB, you can effectively use about 60GB as the OS uses memory

Priorities of Jobs

In order of importance, jobs are prioritised in this order:

1. Resources available to the project
2. Walltime
3. Number of jobs (fair share)

Best strategy

1. Submit jobs constantly/daily
2. Have about 10-20 jobs in the machine
3. Be realistic with walltime
4. Don't ask for resources you don't need!

NCI Facilities

| Disk Type | Disk Usage |
|------------|---|
| NCI Raijin | <ul style="list-style-type: none">• 2 sockets with 8-core CPU's = 16 cores• 57,472 cores in the compute nodes• Approximately 160 TBytes of main memory;• Infiniband FDR (fourteen data rate) interconnect• Approximately 10 PBytes of usable fast file system (for short-term scratch space). |

NCI Facilities

Raijin

96.5% of Nodes have 24Gb/node

3.2% of Nodes have 48Gb/node

0.3% of Nodes have 96Gb/node

Software on NCI

| Area | Software |
|--------------------------|---|
| Computational Chemistry | ABINIT, Amber, CPMD*, GULP*, NAMD*, Molpro etc. |
| Bioinformatics | AbySS, BEAST, BIOPERL, Cufflinks, MAW, etc. |
| Math Libraries | ARPACK, BLACS, Boost, FFTW, GSL, MKL, Tao |
| Statistics & Maths Env's | Maple*, Mathematica*, MatLab*, Octave*, R, Stata* |

- Asterisked items indicates that discussion with NCI facility staff is required before use (Licensing issues)
- <http://nf.nci.org.au/facilities/software/index.php>

Software on NCI

| Area | Software |
|--------------------------|---|
| Computational Chemistry | ABINIT, Amber, CPMD*, GULP*, NAMD*, Molpro etc. |
| Bioinformatics | AbySS, BEAST, BIOPERL, Cufflinks, MAW, etc. |
| Math Libraries | ARPACK, BLACS, Boost, FFTW, GSL, MKL, Tao |
| Statistics & Maths Env's | Maple*, Mathematica*, MatLab*, Octave*, R, Stata* |

- Asterisked items indicates that discussion with NCI facility staff is required before use (Licensing issues)
- <http://nf.nci.org.au/facilities/software/index.php>

Orange Physical Disks

- Which of the 3 disks to use and when?

| Type | Disk Type | Disk Type | • Disk Usage |
|------|-----------|-----------|---|
| 1 | Panasas | 59Tb | <ul style="list-style-type: none">• Parallel global file system• All nodes see the Panasas disks<ul style="list-style-type: none">• directly attached to each node• Very fast for large files• Can be slow if you copy MANY small files• System director blade creates metadata for each file |
| 2 | SGI | 50Tb | <ul style="list-style-type: none">• An NFS mounted file system• All nodes see the SGI disks<ul style="list-style-type: none">• can only be seen via NFS backbone• Uses old technology, therefore very robust• Scales nicely for clusters up to 100 Nodes (very good for Orange) |

Orange Disks (cont.)

- Which of the 3 disks to use and when?

| Type | Disk Type | Disk Type | • Disk Usage |
|------|---------------|-----------|--|
| 3 | Local Scratch | 200Tb | <ul style="list-style-type: none">• <u>Exist in each node</u>• No network is necessary making these the fastest disk• If you have a lot of I/O, you should copy your data to here and work here• One 2Tb disk in each compute node• Is only accessible within that node as it's not attached to any network, therefore cannot be accessed by another node• <u>NOTE:</u> The lifetime of files on this disk is only for the duration of the runtime of the job – The user must copy back results. If not, the next job will erase any files |

Disk Partitions - Orange

| /home | |
|----------------|-------------------------------|
| Mounted under: | /home/username |
| Disk Type | SGI Disks |
| Size: | 60GB default |
| Backed up: | Yes |
| Speed: | Intermediate disk (SGI Disks) |
| Life time: | Permanent |

Disk Partitions - Orange

| /projects/project-name | |
|-------------------------------|--|
| Mounted under: | /projects/project-name |
| Disk Type | Panasas Disk |
| Size: | no default size |
| Backed up: | Yes |
| Speed: | High speed |
| Life time: | Till end of the running year - merit allocation period |

- There will also be some “repository space” for large datasets, such as bioinformatics databases

Disk Partitions - Orange

| /data2 | |
|----------------|---------------------|
| Mounted under: | /data2 on each node |
| Disk Type | Scratch Disks |
| Size: | Limit of disk - 2TB |
| Backed up: | No |
| Speed: | Fastest |
| Life time: | Job duration |

Warning: This partition is shared among users, so can be “filled up” (with other jobs) while your job is running!

Disk Partitions

You can find out more about the partitions on the HPC machine using the **df** command.

| Command | Description |
|--------------|--|
| df -h | Show disk free space for all partitions in human readable format |

You can find out more about current disk usage, using the **du** command.

| Command | Description |
|-----------------|---|
| du -hs . | Show disk usage of current directory in human readable format |

Quotas

- There is no quota on scratch disks for performance reasons
- The quota on /projects/project-name depends on your allocation
- 60 GB soft limit for /home
- 80 GB hard limit for /home (30 days)

More Info on NCI & Orange

- Read more about Orange and NCI Facilities
 - <http://www.intersect.org.au/hpc-news>
 - <http://www.intersect.org.au/orange>
 - http://www.intersect.org.au/nci_next
- Sample PBS Script & Info on Orange
 - <http://www.intersect.org.au/orange-handbook>

Resource Allocation Round

- Merit-based system by which Intersect members can gain access to our HPC facilities
- Applications reviewed by HPC staff (for technical complexity and track record) and the Intersect Resource Allocation Committee (for research merit)
- Applications to Intersect's HPC systems will be made through NCI's forms in October each year
- Applications must be made by Academic Staff at an Intersect member institutions. PhD students can make use of the facilities, the lead CI must be an academic staff member.
- Questions to: hpc_support@intersect.org.au

Register with NCI (step 1)

Register **a new Id with NCI:**

[http://nf.nci.org.au/accounts/forms/user registration.php](http://nf.nci.org.au/accounts/forms/user_registration.php)

This will provide your details to NCI

Register with NCI (step 2)

Apply for **a project from NCI:**

[https://nf.nci.org.au/accounts/projects_new/
APP_form.php](https://nf.nci.org.au/accounts/projects_new/APP_form.php)

This will provide link your Id to your Project

Project Registration Form

- Pick INTERSECT under partner/scheme on the first page of the project registration form or else you won't get access to Orange
- If you're unsure about which machine to get access to, email hpc_support@intersect.org.au who can advise you
- You can add accounts to an existing project also!

Conclusion

- In this course we have covered
 - the basics of the Unix command line
 - transferring data
 - the specifics of our HPC machine
- As different machines have different PBS systems, scripts that work on Octane may not work on NCI facilities

Thanks for attending!

- Please complete our course survey at:
 - <http://svy.mk/18c8dHa>
- Any further questions, please contact us at training@intersect.org.au