HPC

Introduction to Unix and HPC

INTERSECT

What is HPC?

- HPC, or high-performance computing, refers to the application of supercomputers or clusters of computers to computational problems that typically arise through scientific inquiry
- HPC is useful when a computational problem:
  - **<u>Is too large</u>** to solve on a conventional laptop or workstation (because it requires too much memory or disk space) or
  - **<u>Would take too long</u>** (because the algorithm is complex, the dataset is large, or data access is slow) or
  - **<u>Are too many</u>** – High Throughput Computing

Parallelism on HPC

- HPC systems often derive their computational power by <u>exploiting parallelism</u>
- Programs for HPC systems must be split up into many smaller "sub-programs" which can be executed in parallel on different processors
- HPC systems can offer parallelism at a much larger scale, with 100's or 1000's, or (soon) even <u>millions of tasks running concurrently</u>.
- Writing <u>parallel software can be challenging</u>, and many existing software packages do not already support parallelism & may require development.
  <u>NOTE: Many tasks cannot be parallelised</u>

Reasons to use HPC

- You have a program that can be recompiled or reconfigured to use optimized numerical libraries that are available on HPC systems but not on your own system.
- <u>HPC applications are already installed </u>on the HPC machines which is a non-trivial task
- <u>You have a "parallel" problem</u>, e.g. you have a single application that needs to be rerun many times with different parameters.
- You have an application that has already been <u>designed with parallelism</u>
- To make use of the <u>large memory </u>available
- Our facilities are <u>reliable </u>and regularly backed up

When not to use HPC?

- You have a single threaded job which will only run one job at a time (typical of MatLab users)
- You rely on Databases
- You have a lot of data to transfer between your local machine and the HPC on a continuous basis (e.g. per job)
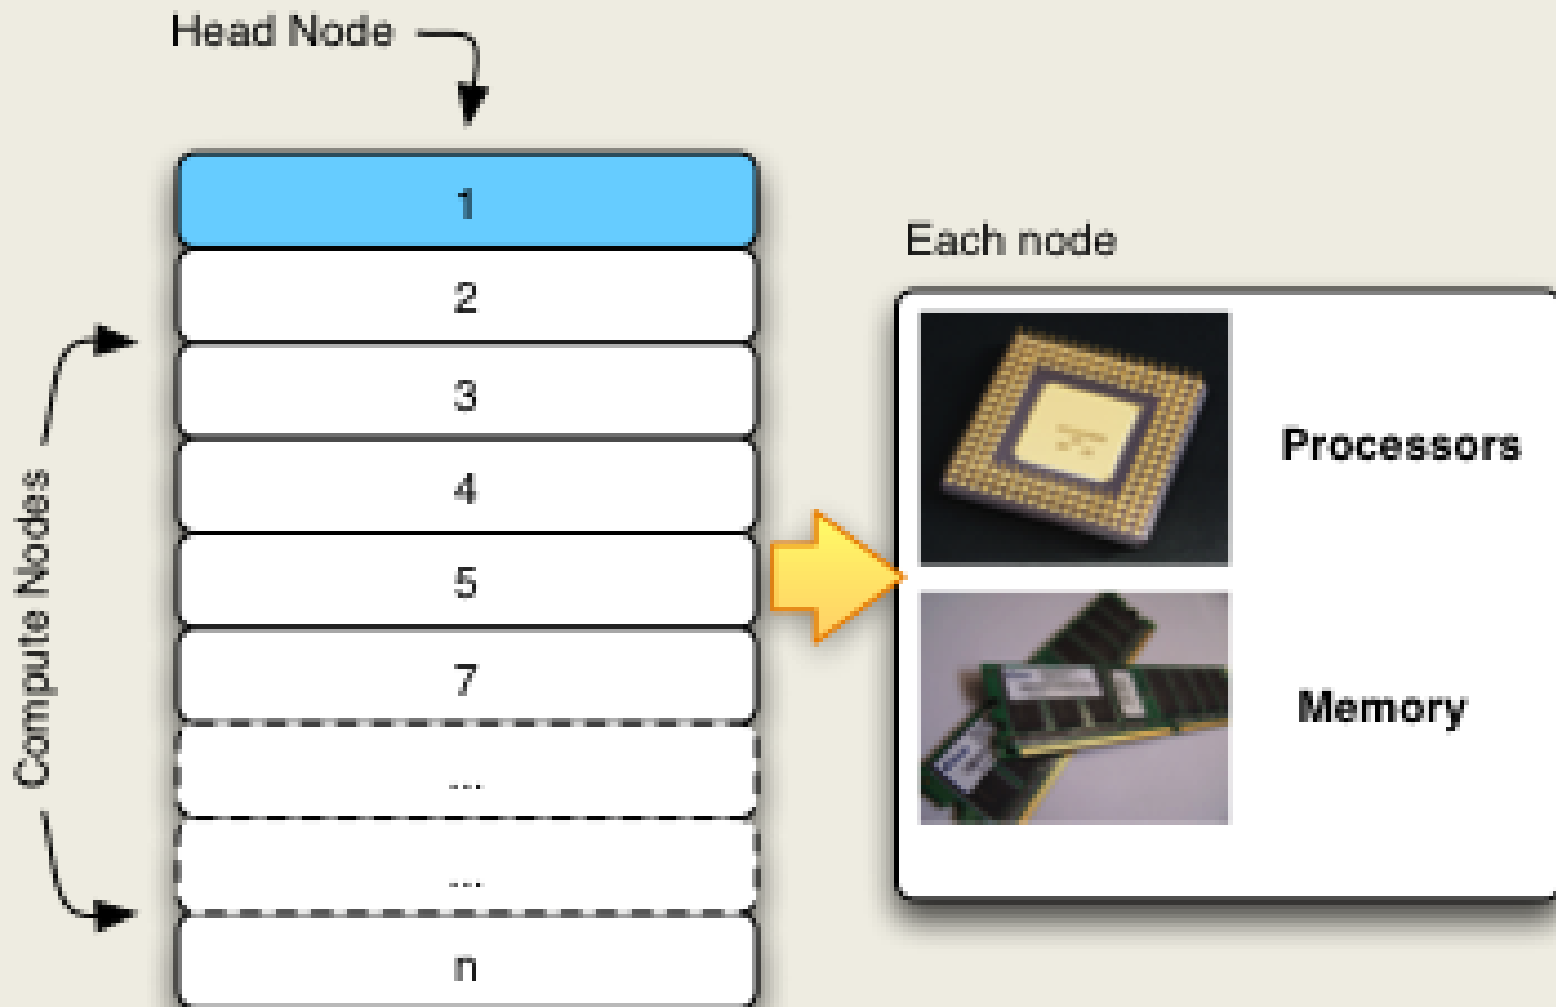- You need to have a GUI to interact with your program

HPC machines

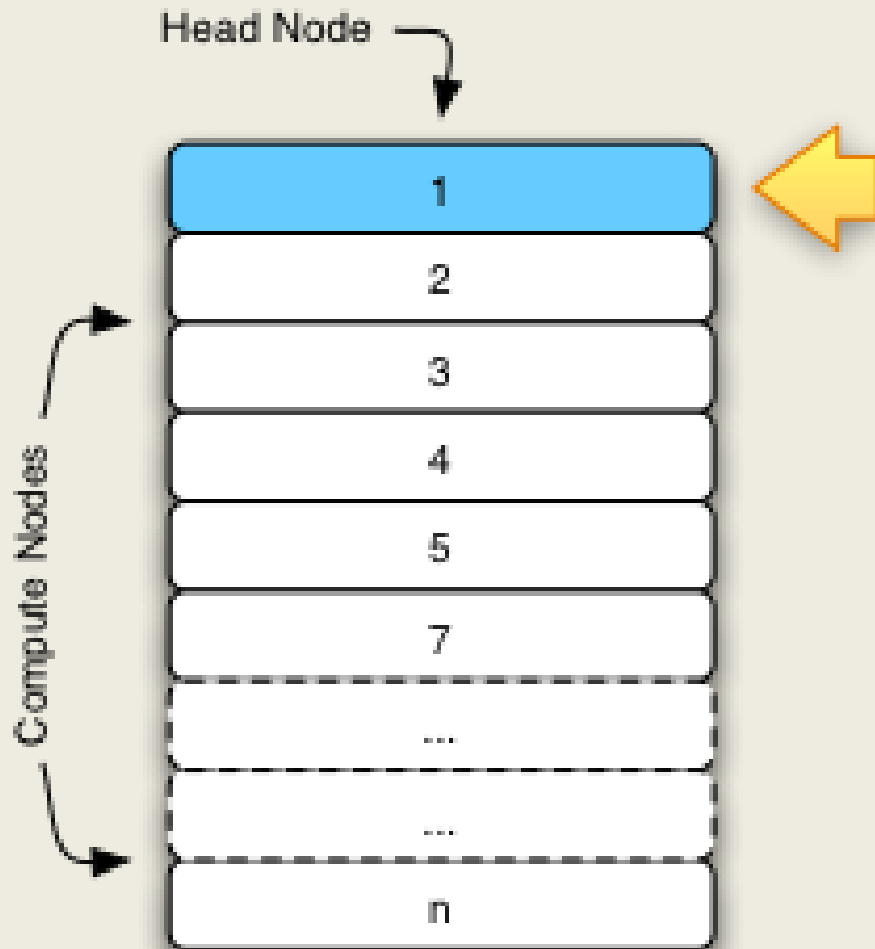| System	Memory Architecture	Cores	Nodes	Memory
Octane (training machine)	Distributed	48	3	48GB
Orange	**Distributed**	**1,600**	**100**	**8TB**
NCI – (Vayu)	**Distributed**	**11,936**	**1492**	**37TB**
NCI – (Name TBA) • **available May 2013**	**Distributed**	**57,472**	**3592**	**158TB**

The typical HPC workflow

- In HPC we talk about **jobs**, these are simply commands we wish to run and requests for resources (e.g. compute time, disk space, memory requirements, setup of s/w env's etc.)
- Generally time consuming & resource intensive.
- Jobs are typically run **non-interactively**,
- Can be run **interactively** for testing purposes
- We add our jobs to a **queue**.
- When machines have **free resources** jobs run
- Once jobs are complete, we can inspect their output.

The HPC "Cluster"

Head Node

Compute Nodes

1
2
3
4
5
7
...
...
n

Each node

Processors

Memory

The Login Node

Head Node

Compute Nodes

1
2
3
4
5
7
...
...
n

Login Node
- Interactive programs
- SSH sessions
- Testing
- Compiling

The Head Node

Head Node
- Queuing jobs

Compute Nodes

- These run your jobs
- Managed by the **scheduler**
- Typically you will not interact with the nodes directly (some users may need to)

Queuing Systems

- **Portable Batch System** (PBS) is the name of computer software that performs job scheduling. Its primary task is to <u>allocate computational tasks, i.e., batch jobs, among the available computing resources</u>.
- The following versions of PBS are currently available:
  - OpenPBS
  - TORQUE
  - PBS Professional (PBS Pro)
  - ANU PBS
- Guide to PBS: http://hpc.sissa.it/pbs/pbs.html

Queuing Systems cont.

- Another popular batch system is **SLURM** (Simple Linux Utility for Resource Management)
  - Open source, fault-tolerant, and highly scalable cluster management and job scheduling system for large and small Linux clusters.
  - Very useful for use on clusters
  - Platform Tools used by IBM
  - Used by many supercomputers, e.g. TERA 100 at CEA (Europe's most powerful supercomp.)
- Many banks and commercial entities using batch systems

ANU PBS vs PBS Pro

- ANU PBS is a **<u>customised version of PBS based on OpenPBS 2.3</u>** maintained by ANU
- Details of ANU PBS modifications are found here: http://anusf.anu.edu.au/~dbs900/PBS/local_modifications.html

Batch System	ANU PBS	PBS PRO
Machines using	Vayu & NCI New Facility	Orange & Octane
Code Base	OpenPBS 2.3	PBS Professional
Licence	ANU Licence	Altair Licence

The Batch Queuing System

Batch Queuing component

Head Node

Batch Queue Commands

- The batch system is a normal program
- Lets you add and remove jobs from the queue and monitor the queue
- Script/command line driven

The Scheduler component

- Allocates jobs to compute nodes
- Optimizes usage of resources
- "Optimize" can mean many things
- Non-trivial
- Never interact with directly

Batch Queue

Compute Nodes

Pluck

Run

Scheduler

PBS Pro Commands

In order to use the batch system productively, we need to know how to perform three actions:

- Add a job to the queue
- Remove a job from the queue
- See where our job is in the queue

Command	Description
`qsub <job-script>`	Submit a job (add to queue) Returns a <job-number>
`qdel <job-number>`	Delete job (remove from the queue)
`qstat <job-number>`	Monitor jobs
`qalter <job-number>`	Modifies the attributes of the job or jobs

Exercise 1

Monitoring the queue with qstat

Command	Description
`qstat -a`	List all jobs in the queue
`qstat -u <username>`	List all jobs of a particular user
`qstat -f <job-number>`	Show detailed information about a job
`nqstat`	ANU PBS only! Shows detailed information about all jobs

INTERSECT

All about Modules

- **What do Modules do?**
  - Set up the environment for a software package
  - Adds paths to executables to $PATH
  - May change other shell variables, and/or load other modules.
  - Allow you to have different versions of the same software package, e.g. load the Intel compilers v23 with module load intel/12, or if you require an older version, you can load it with module load intel/9

All about Modules (cont.)

Things to know about modules
- No <u>default module loaded when you login</u>. You can change this by adding lines to your .bashrc file (BASH users) or similar for other Shells
- Modules exist for many packages but <u>not for all</u>!
- <u>Some modules exclude each other</u>, e.g.
  - You can't load Intel compilers v8 & v9. You can only load one.
  - You can only load 1 MPI (e.g. MPT or Intel or OpenMPI)

Module Commands

Command	Description
`module avail`	Will list all available module files in the current MODULEPATH
`module load/unload`	Will load/unload a modulefile into the shell environment e.g. `module load mpt/2.06`
`module list`	List all loaded modules
`module show [modulefile]`	Will show information about the modulefile

Add a job to the queue

- To add a job to the queue, we write a **job script**. A job script is a simply a script.
- The # symbol signifies a comment for the Shell
- It has some special comments that pass info to PBS Pro.
- **#PBS** is a keyword for PBS and specifies that this line is for PBS. The Shell will ignore it
- When we want to queue the job, we pass its filename as a parameter to **qsub**, e.g.
  - **qsub <job-script>**
- The batch queuing system will return a number that uniquely identifies the job.

Useful Environment Variables

- These are available in the context of your job script.

Command	Description
PBS_O_WORKDIR	The directory the job was submitted from
PBS_JOBID	The job number given when the job was submitted

A sample PBS job script

NOTE: This script will not run on Orange as resources are assigned via Groups on Orange, not via projects

```bash
#!/bin/bash
# Request resources
# * 10 minutes wall time to run
#PBS -l walltime=00:10:00
# * 1 node, 1 processor
#PBS -l nodes=1:ppn=1
# * 100 megabytes physical memory allocated to job
#PBS -l mem=100mb
# Specify a project code (for accounting)
#PBS -P a40
cd $PBS_O_WORKDIR
# Specify the job to be done
date
sleep 10
date
```

You've got mail!

```
# Set email address
#PBS -M fred@intersect.org.au
# Send an email when jobs
# begins (b), gets aborted (a)
# and ends (e)
#PBS -m abe
```

Exercise 2

Create a script and submit a sample job

Command	Description
`#PBS`	`#PBS` is a keyword for PBS and specifies that this line is for PBS. The Shell will ignore it
`#`	Signifies a comment for the Shell, e.g. `# Next line will create a job`
`qsub`	Submit a job to PBS
`qstat`	List jobs in the queue
`cat <filename>`	Print a file to the terminal (catenate)
`less <filename>`	Like `cat`, but less at a time
`nano <filename>`	Will open file *<filename>* in a text file editor
`Sample PBS Script`	http://www.intersect.org.au/orange-handbook

INTERSECT

Job limits on Orange

- 200 hours of **walltime**
- 64GB of **memory** per standard node. e.g. 128GB for 2 nodes etc.
- 256GB of **memory** per large memory nodes
- **NOTE**: If you grab a node with 64GB, you can effectively use about 60GB as the OS uses memory

Priorities of Jobs

In order of importance, jobs are prioritised in this order:

1. Resources available to the project

2. Walltime

3. Number of jobs (fair share)

Best strategy

- Submit jobs constantly/daily

- Have about 10-20 jobs in the machine

- Be realistic with walltime

- Don't ask for resources you don't need!

NCI Facilities

Disk Type	Disk Usage
VAYU **(Commissioned in 2010)**	• Sun Constellation Cluster with 1492 nodes, each containing 2 quad core Nehalem processors summing up to 11,936 cores. 37TB RAM and 800 TB disk space. Uses **QDR** (quad data rate) – predecessor to FDR • Global memory address space where shared memory programs can run • The unit of **shared memory parallelism is the node**, comprised of **2 sockets with 4-core CPU's = 8 cores**.
NCI Upcoming System **(available in May)**	• **2 sockets with 8-core CPU's = 16 cores** • 57,472 cores in the compute nodes • Approximately 160 TBytes of main memory; • Infiniband **FDR** (fourteen data rate) interconnect • Approximately 10 PBytes of usable fast file system (for short-term scratch space).

NCI Facilities

VAYU	New NCI Facilities
66% of Nodes have 32Gb (2Gb/core)	96.5% of Nodes have 24Gb/node
33% of Nodes have 64Gb (4Gb/core)	3.2% of Nodes have 48Gb/node
2% of Nodes have 128Gb (8Gb/core)	0.3% of Nodes have 96Gb/node

Software on NCI

Area	Software
Computational Chemistry	ABINIT, Amber, CPMD*, GULP*, NAMD*, Molpro etc.
Bioinformatics	AbySS, BEAST, BIOPERL, Cufflinks, MAW, etc.
Math Libraries	ARPACK, BLACS, Boost, FFTW, GSL, MKL, Tao
Statistics & Maths Env's	Maple*, Mathematica*, MatLab*, Octave*, R, Stata*

- Asterisked items indicates that discussion with NCI facility staff is required before use (Licensing issues)
- **http://nf.nci.org.au/facilities/software/index.php**

Orange Physical Disks

- Which of the 3 disks to use and when?

Type	Disk Type	Disk Type	• Disk Usage
1	Panasas	59Tb	• Parallel global file system • All nodes see the Panasas disks • directly attached to each node • Very fast for large files • Can be slow if you copy MANY small files • System director blade creates metadata for each file
2	SGI	50Tb	• An NFS mounted file system • All nodes see the SGI disks • can only be seen via NFS backbone • Uses old technology, therefore very robust • Scales nicely for clusters up to 100 Nodes (very good for Orange)

Orange Disks (cont.)

- Which of the 3 disks to use and when?

Type	Disk Type	Disk Type	• Disk Usage
3	Local Scratch	200Tb	• <u>Exist in each node</u> • No network is necessary making these the fastest disk • If you have a lot of I/O, you should copy your data to here and work here • One 2Tb disk in each compute node • Is only accessible within that node as it's not attached to any network, therefore cannot be accessed by another node • **<u>NOTE:</u>** The lifetime of files on this disk is only for the duration of the runtime of the job – The user must copy back results. If not, the next job will erase any files

Disk Partitions - Orange

/home	
Mounted under:	/home/username
Disk Type	SGI Disks
Size:	60GB default
Backed up:	Yes
Speed:	Intermediate disk (SGI Disks)
Life time:	Permanent

Disk Partitions - Orange

/projects/project-name	
Mounted under:	/projects/project-name
Disk Type	Panasas Disk
Size:	no default size
Backed up:	Yes
Speed:	High speed
Life time:	Till end of the running year - merit allocation period

- There will also be some "repository space" for large datasets, such as bioinformatics databases

Disk Partitions - Orange

/data2	
Mounted under:	/data2 on each node
Disk Type	Scratch Disks
Size:	Limit of disk - 2TB
Backed up:	No
Speed:	Fastest
Life time:	Job duration

Warning: This partition is shared among users, so can be "filled up" (with other jobs) while your job is running!

Disk Partitions

You can find out more about the partitions on the HPC machine using the **df** command.

Command	Description
df -h	Show disk free space for all partitions in human readable format

You can find out more about current disk usage, using the **du** command.

Command	Description
du -hs .	Show disk usage of current directory in human readable format

Quotas

- There is no quota on scratch disks for performance reasons.

- The quota on /projects/project-name depends on your allocation.

- 60 GB soft limit for /home.
- 80 GB hard limit for /home (30 days).

More Info on NCI & Orange

- Read more about Orange and NCI Facilities
  - http://www.intersect.org.au/hpc-news
  - http://www.intersect.org.au/orange
  - http://www.intersect.org.au/nci_next
- Sample PBS Script & Info on Orange
  - http://www.intersect.org.au/orange-handbook

Resource Allocation Round

- <u>Merit-based system</u> by which Intersect members can gain access to our HPC facilities
- Applications reviewed by HPC staff (for <u>technical complexity</u> and track record) and the Intersect Resource Allocation Committee (for <u>research merit</u>)
- Applications to Intersect's HPC systems will be made through <u>NCI's forms in October each year</u>
- Applications must be <u>made by Academic Staff</u> at an Intersect member institutions. PhD students can make use of the facilities, the lead CI must be an academic staff member.
- Questions to: hpc_support@intersect.org.au

Register with NCI (step 1)

Register **a new Id with NCI**:

http://nf.nci.org.au/accounts/forms/user_registration.php

This will provide your details to NCI

Register with NCI (step 2)

Apply for **<u>a project from NCI</u>**:

https://nf.nci.org.au/accounts/projects_new/APP_form.php

<u>This will provide link your Id to your Project</u>

Project Registration Form

- Pick <u>INTERSECT under partner/scheme</u> on the first page of the project registration form or else you won't get access to Orange
- If you're unsure about which machine to get access to, email hpc_support@intersect.org.au who can advise you
- You can add accounts to an existing project also!

Conclusion

- In this course we have covered the basics of the Unix command line, transferring data, and the specifics of our HPC machine
- As different machines have different PBS systems, scripts that work on Octane may not work on NCI facilities
- Please complete our survey!
- Any questions?